

## Отчет

1. Добавляем контейнеры zookeeper, kafka и consumer в docker-compose.yml файл. Для брокера прописываем порты, адрес доступа, количество партиций и репликаций топика. Добавляем healthcheck для того, чтобы другие контейнеры стартовали после запуска kafka контейнера. В контейнер consumer выносим запуск предобработки данных и тренировки, поскольку именно он будет отвечать за работу с данными приходящими по api. Последним шагом запускаем сам консьюмер и создаем файл consumer-ready в качестве флага для начала работы web контейнера. Web сервис теперь просто отвечает за запуск приложения.

### Листинг 1 - Код docker-compose.yml.

```
services:
  decrypt:
    image: python:3.9-slim
    container_name: decrypt_env
    volumes:
      - ./env:/app/env
    working_dir: /app/env
    command: bash -c "
      ls -la
      && source vault-pass.env
      && pip install ansible-vault
      && echo $$VAULT_PASSWORD > vault-pass.txt
      &&          ansible-vault          decrypt          secrets.env
--vault-password-file vault-pass.txt --output .env
      && rm -rf vault-pass.txt"
  redis:
    image: redis:latest
    container_name: redis_db
    depends_on:
      decrypt:
        condition: service_completed_successfully
```

```
volumes:
  - ./env:/app/env
working_dir: /app/env
command: bash -c "
  ls -la
  && source .env
  && redis-server --requirepass $$REDIS_PASSWORD"
ports:
  - 6379:6379
healthcheck:
  test: ["CMD", "bash", "-c", "
    cd /app/env
    && source .env
    && redis-cli -a $$REDIS_PASSWORD ping"]
  interval: 1s
  timeout: 2s
  retries: 10
zookeeper:
image: confluentinc/cp-zookeeper:7.3.0
container_name: zookeeper
environment:
  ZOOKEEPER_CLIENT_PORT: 2181
  ZOOKEEPER_TICK_TIME: 2000
healthcheck:
  test: [ "CMD", "nc", "-vz", "localhost", "2181" ]
  interval: 10s
  timeout: 3s
  retries: 3
broker:
image: confluentinc/cp-kafka:7.3.0
hostname: broker
container_name: broker
depends_on:
  zookeeper:
    condition: service_healthy
environment:
  KAFKA_BROKER_ID: 1
```

```

    KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
    KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://broker:9092
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
    KAFKA_NUM_PARTITIONS: 1
healthcheck:
    test: ["CMD", "kafka-broker-api-versions",
"--bootstrap-server", "broker:9092"]
    interval: 10s
    timeout: 10s
    retries: 5
ports:
    - 9092:9092
consumer:
build: .
container_name: consumer
volumes:
    - ./env:/app/env
command: bash -c "
    ls -la
    && python src/preprocess.py
    && python src/train.py
    && python src/predict.py -m LOG_REG -t func
    && coverage run src/unit_tests/test_preprocess.py
    && coverage run -a src/unit_tests/test_training.py
    && coverage run -a src/unit_tests/test_database.py
    && coverage report -m
    && (python src/kafka_consumer.py --model LOG_REG &)
    && touch /tmp/consumer-ready
    && tail -f /dev/null"
depends_on:
    decrypt:
        condition: service_completed_successfully
    redis:
        condition: service_healthy
    broker:
        condition: service_healthy
healthcheck:

```

```
test: ["CMD-SHELL", "test -f /tmp/consumer-ready"]
interval: 1m
timeout: 10s
start_period: 5m
retries: 30
web:
build: .
container_name: web
volumes:
    - ./env:/app/env
command: bash -c "
    ls -la
    && coverage run -a src/unit_tests/test_database.py
    && coverage run -a src/unit_tests/test_app.py
    && coverage report -m
    && (python src/app.py &)
    && sleep 30
    && curl -X GET http://localhost:8000/
    && curl -X POST http://localhost:8000/predict \
        -H 'Content-Type':" application/json' \
        --data-binary @tests/test_0.json"
depends_on:
    decrypt:
        condition: service_completed_successfully
    redis:
        condition: service_healthy
    broker:
        condition: service_healthy
    consumer:
        condition: service_healthy
ports:
    - 8000:8000
image: zarus03/ml-big-data-lab-4:latest
```

2. Выделим консьюмер в отдельный класс `kafka_consumer.py`. В нем запускается специальная модель, и в цикле обрабатываются приходящие с продюсера сообщения.

#### Листинг 2 - Реализация `kafka_consumer.py`.

```
import json
import pandas as pd
import pickle
import traceback
import sys
import argparse
import configparser
from logger import Logger
from database import RedisClient
from confluent_kafka import Consumer, KafkaError

SHOW_LOG = True

class KafkaConsumer:
    def __init__(
        self,
        args: argparse.Namespace,
        broker: str,
        topic: str,
        group: str
    ):
        logger = Logger(SHOW_LOG)
        self.log = logger.get_logger(__name__)

        self.config = configparser.ConfigParser()
        self.config.read("config.ini")

        self.args = args

        self.model, self.scaler = self._load_model()
```

```

self.log.info('Kafka Consumer Model initialized')

self.database_client = RedisClient()

conf = {
    'bootstrap.servers': broker,
    'group.id': group,
    'session.timeout.ms': 6000,
    'auto.offset.reset': 'earliest',
    'enable.auto.offset.store': False}

self.topic = topic

self.consumer = Consumer(conf)
self.consumer.subscribe(topics=[self.topic])

def _load_model(self):
    """
    Loads the machine learning model and scaler from disk.

    Returns:
        tuple: Loaded model and scaler.
    """
    try:
        with open(self.config[self.args.model]["path"], "rb") as
model_file:
            model = pickle.load(model_file)

        with open(self.config["STD_SCALER"]["path"], "rb") as
scaler_file:
            scaler = pickle.load(scaler_file)

        return model, scaler
    except FileNotFoundError:
        self.log.error(traceback.format_exc())
        sys.exit(1)

```

```

def process(self, data: dict):
    prediction_id = data['prediction_id']
    input_data = json.loads(data['input_data'])

    X = self.scaler.transform(pd.json_normalize(input_data['X']))
    pred = self.model.predict(X).tolist()

    prediction_data = {'prediction': pred}

    self.database_client.set(prediction_id,
                              json.dumps(prediction_data))

def consume_messages(self, timeout: float = 1.0):
    try:
        self.log.info(f"Starting consumer for topic {self.topic}")
        while True:
            msg = self.consumer.poll(timeout)

            if msg is None:
                self.log.debug("No message received")
                continue

            if msg.error():
                if msg.error().code() ==
KafkaError._PARTITION_EOF:
                    self.log.error(f'Reached end of partition:
{msg.topic()} [{msg.partition()}]')
                else:
                    self.log.error(f'Error while consuming
messages: {msg.error()}')
                    continue

            self.log.info(f"Received message on consumer one:
{msg.value().decode('utf-8')}")

```

```

self.process(json.loads(msg.value().decode('utf-8')))

    except KeyboardInterrupt:
        self.log.info("Stopping consumer...")

    finally:
        self.consumer.close()
        self.log.info("Consumer closed")

def main(args):
    consumer = KafkaConsumer(
        args,
        broker="broker:9092",
        topic="predictions",
        group="ml_app_group"
    )
    consumer.consume_messages()

if __name__ == "__main__":
    parser = argparse.ArgumentParser(description="Kafka
Consumer")
    parser.add_argument("-m", "--model",
                        type=str,
                        help="Select model",
                        required=True,
                        default="LOG_REG",
                        const="LOG_REG",
                        nargs="?",
                        choices=["LOG_REG"])
    args = parser.parse_args()
    main(args)

```



3. Выделим продюсер тоже в отдельный класс `kafka_producer.py`. В нем реализован специальный метод по отправке сообщения в брокер.

### Листинг 3 - Реализация `kafka_consumer.py`.

```
import json
from logger import Logger
from confluent_kafka import Producer

SHOW_LOG = True

class KafkaProducer:
    def __init__(
        self,
        broker: str,
        topic: str
    ):
        logger = Logger(SHOW_LOG)
        self.log = logger.get_logger(__name__)

        producer_config = {
            'bootstrap.servers': broker,
        }
        self.topic = topic

        self.producer = Producer(producer_config)

    def delivery_report(self, err, msg):
        if err is not None:
            self.log.error(f'Message delivery failed: {err}')
        else:
            self.log.info(f'Message delivered to {msg.topic()}
[ {msg.partition()} ]')

    def send_message(self, message: dict):
```

```

try:
    data = json.dumps(message).encode('utf-8')
    self.producer.produce(
        self.topic,
        value=data,
        callback=self.delivery_report
    )
    self.producer.flush()

except Exception as e:
    self.log.error(f'Failed to send message: {e}')
    raise

```

4. Сам продюсер запускается в коде приложения. Отправки сообщения в брокер происходит в момент получения запроса на предсказание по ручке /predict.

#### Листинг 4 - Реализация app.py.

```

from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
from typing import List, Dict
from logger import Logger
from database import RedisClient
from kafka_producer import KafkaProducer
import traceback
import uvicorn
import json
import uuid

SHOW_LOG = True

class PredictionInput(BaseModel):
    """
    Pydantic model for input data validation.

```

```

Attributes:
X (List[Dict[str, float]]): List of feature dictionaries.
y (List[Dict[str, float]]): List of target dictionaries.
"""

X: List[Dict[str, float]]
y: List[Dict[str, float]]

class WebApp:
    """
    Web application class using FastAPI for serving a machine
    learning model.
    """

    def __init__(self):
        """
        Initializes the web application, loads the model, and creates
        the FastAPI app.

        Args:
            args (argparse.Namespace): Parsed command-line
            arguments.
        """
        logger = Logger(SHOW_LOG)
        self.log = logger.get_logger(__name__)

        self.app = self._create_app()
        self.log.info('FastAPI app initialized')

        self.database_client = RedisClient()

        self.kafka_producer = KafkaProducer(
            broker="broker:9092",
            topic="predictions",
        )

    def _create_app(self):
        """

```

Creates and configures the FastAPI application.

Returns:

FastAPI: Configured FastAPI app instance.

"""

```
app = FastAPI()
```

```
@app.get("/")
```

```
async def root():
```

```
    """Root endpoint for health check."""
```

```
    return {"message": "Hello World"}
```

```
@app.post("/predict")
```

```
async def predict(input_data: PredictionInput):
```

```
    """
```

```
    Endpoint for making predictions using the trained model.
```

Args:

input\_data (PredictionInput): Input data containing X (features) and y (target values).

Returns:

dict: A dictionary containing predictions and model score.

```
    """
```

```
    try:
```

```
        prediction_id = str(uuid.uuid4())
```

```
        prediction_data = {
```

```
            "prediction_id": prediction_id,
```

```
            "input_data": input_data.model_dump_json()
```

```
        }
```

```
        self.kafka_producer.send_message(prediction_data)
```

```
        return {"prediction_id": prediction_id}
```

```

        except Exception as e:
            self.log.error(traceback.format_exc())
            raise HTTPException(status_code=500,
detail=str(e))

@app.get("/predictions/{prediction_id}")
async def get_prediction(prediction_id: str):
    """
    Endpoint for retrieving a prediction from Redis by its
ID.

    Args:
        prediction_id (str): The ID of the prediction to
retrieve.

    Returns:
        dict: The prediction data stored in Redis.
    """
    prediction_data = self.database_client.get(prediction_id)
    if prediction_data is None:
        raise HTTPException(status_code=404,
detail="Prediction not found")
    return json.loads(prediction_data)

return app

def run(self, host: str = "0.0.0.0", port: int = 8000):
    """
    Runs the FastAPI application using Uvicorn.

    Args:
        host (str): Host address to run the server.
        port (int): Port number to run the server.
    """
    uvicorn.run(self.app, host=host, port=port)

```

```
if __name__ == "__main__":  
    web_app = WebApp()  
    web_app.run()
```